

# CS 305 Computer Networks

## Chapter 2 Application Layer (I)

Jin Zhang

Department of Computer Science and Engineering  
Southern University of Science and Technology

# Chapter 2: outline

## 2.1 principles of network applications

## 2.2 Web and HTTP

## 2.3 electronic mail

- SMTP, POP3, IMAP

## 2.4 DNS

## 2.5 P2P applications

## 2.6 video streaming and content distribution networks

## 2.7 socket programming with UDP and TCP

# Chapter 2: application layer

## our goals:

- conceptual, implementation aspects of network application protocols
  - transport-layer service models
  - client-server paradigm
  - peer-to-peer paradigm
  - content distribution networks
- learn about protocols by examining popular application-level protocols
  - HTTP
  - FTP
  - SMTP / POP3 / IMAP
  - DNS
- creating network applications
  - socket API

# Some network apps

- e-mail
- web
- text messaging
- remote login
- P2P file sharing
- multi-user network games
- streaming stored video (YouTube, Hulu, Netflix)
- DNS
- voice over IP (e.g., Skype)
- real-time video conferencing
- social networking
- search
- ...
- ...

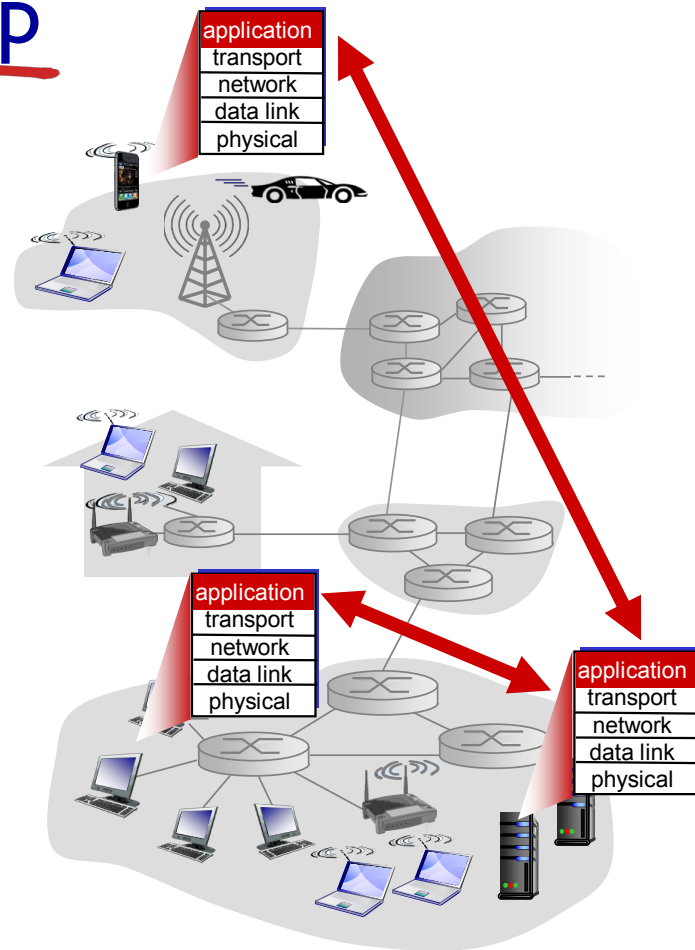
# Creating a network app

write programs that:

- run on (different) end systems
- communicate over network
- e.g., web server software communicates with browser software

no need to write software  
for network-core devices

- network-core devices do not run user applications
- applications on end systems allows for rapid app development, propagation




# Application architectures

possible structure of applications:

- client-server
- peer-to-peer (P2P)

# Client-server architecture



## server:


- always-on host
- permanent IP address
- data centers for scaling

## clients:

- communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do not communicate directly with each other

# P2P architecture

- no always-on server
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
- ★ self scalability — new peers bring new service capacity, as well as new service demands
- peers are intermittently connected and change IP addresses
  - complex management





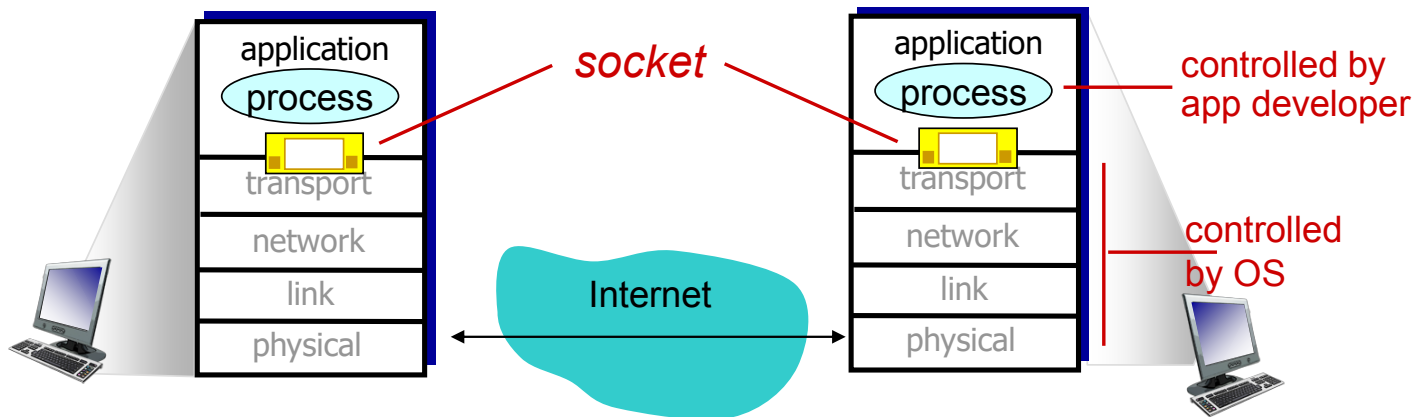
# How to send msg to network?

- Who send/recv msg to/from network? **Processes** (进程)
  - Where does process send/recv msg to/from? **socket**
  - Processes within same host communicate using **inter-process communication** (defined by OS)
  - processes in different hosts communicate by exchanging **messages**
- clients, servers

  - client process**: process that initiates communication
  - server process**: process that waits to be contacted
  - aside: applications with P2P architectures have client processes & server processes

# Sockets

- process sends/receives messages to/from its socket
- socket analogous to door
  - sending process shoves message out door
  - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process



# IP and Port number

- to receive messages, process must have *identifier*
- host device has unique 32-bit IP address
- Q: does IP address of host on which process runs suffice for identifying the process?
  - A: no, many processes can be running on same host (use port to identify)
- *identifier* includes both IP address and port numbers associated with process on host.
- example port numbers:
  - HTTP server: 80
  - mail server: 25
- to send HTTP message to gaia.cs.umass.edu web server:
  - IP address: 128.119.245.12
  - port number: 80

# Requirement What transport service does an app need?

## data integrity 完整性

- some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- other apps (e.g., audio) can tolerate some loss

## timing

- some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

## throughput

- some apps (e.g., multimedia) require minimum amount of throughput to be “effective” *live stream*
- other apps (“elastic apps”) make use of whatever throughput they get

## security

- encryption, data integrity, ...

	Data Integrity	Timing	Throughput.	Security
web	✓	X	X	✓/X
online game	X	✓	X	X
File transfer	✓	X	X	✓/X
online video	X	✓	✓	X

# Transport service requirements: common apps

application	data loss	throughput	time sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video: 10kbps-5Mbps	yes, 100' s msec
stored audio/video	loss-tolerant	same as above	
interactive games	loss-tolerant	few kbps up	yes, few secs
text messaging	no loss	elastic	yes, 100' s msec yes and no

# Internet transport protocols services

## TCP service:

- **reliable transport** between sending and receiving process
- **flow control**: sender won't overwhelm receiver
- **congestion control**: throttle sender when network overloaded
- **does not provide**: timing, minimum throughput guarantee, security
- **connection-oriented**: setup required between client and server processes

## UDP service:

- **unreliable data transfer** between sending and receiving process
- **does not provide**: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup,

Q: why bother? Why is there a UDP?

# Internet apps: application, transport protocols

	<b>application</b>	<b>application layer protocol</b>	<b>underlying transport protocol</b>
	e-mail	SMTP [RFC 2821]	TCP
remote terminal access		Telnet [RFC 854]	TCP
	Web	HTTP [RFC 2616]	TCP
	file transfer	FTP [RFC 959]	TCP
streaming multimedia		HTTP (e.g., YouTube), RTP [RFC 1889]	TCP or UDP
Internet telephony		SIP, RTP, <u>proprietary</u> (e.g., Skype)	TCP or UDP



# App-layer protocol defines

- types of messages exchanged,
  - e.g., request, response
- message syntax: 格式
  - what fields in messages & how fields are delineated
- message semantics 语义
  - meaning of information in fields
- rules for when and how processes send & respond to messages

## open protocols:

- defined in RFCs
- allows for interoperability
- e.g., HTTP, SMTP

## proprietary protocols:

- e.g., Skype

# Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

# Web and HTTP

- *web page* consists of *objects*
- object can be HTML file, JPEG image, Java applet, audio file,...
- web page consists of base HTML-file which includes several referenced objects
- each object is addressable by a *URL*, e.g.,

`www.sustc.edu.cn/resources/cn/image/p27.png`

host name

path name

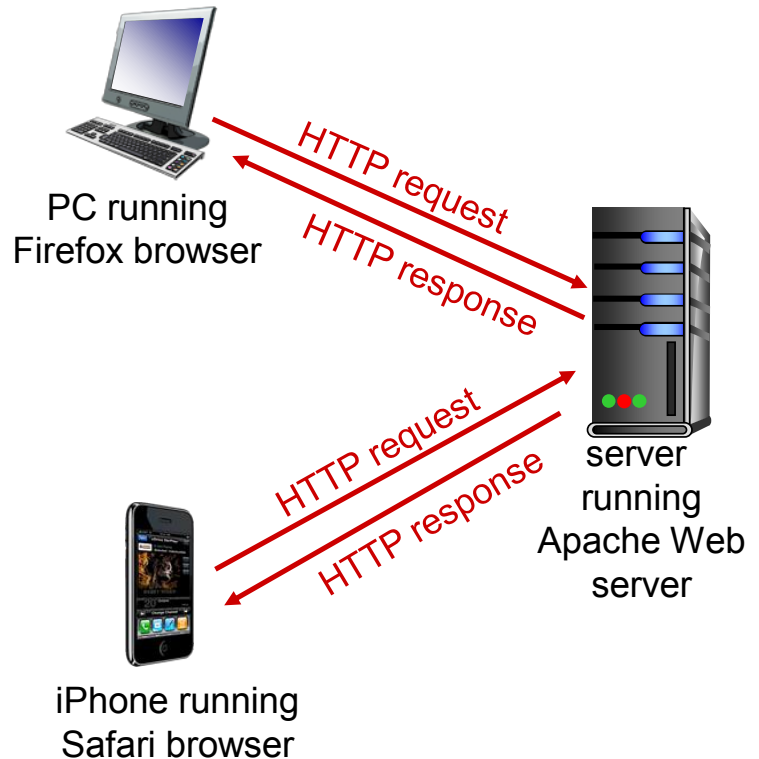
**HTML:** hypertext markup language

**HTTP:** hypertext transfer protocol

# HTTP overview

## HTTP: hypertext transfer protocol

- Web's application layer protocol
- client/server mode
  - **client:** browser that requests, receives, (using HTTP protocol) and "displays" Web objects
  - **server:** Web server sends (using HTTP protocol) objects in response to requests



# HTTP overview (continued)

## *uses TCP:*

- client initiates TCP connection (creates socket) to server, **port 80**
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

## *HTTP is “stateless”*

- server maintains no information about past client requests
- use cookies to check history.*

*aside*

## protocols that maintain “state” are complex!

- past history (state) must be maintained
- if server/client crashes, their views of “state” may be inconsistent, must be reconciled

# HTTP connections

## *non-persistent HTTP*

- at most one object sent over TCP connection
  - connection then closed
- downloading multiple objects required multiple connections

$$(1+1)RTT + 1t_0$$

## *persistent HTTP*

- multiple objects can be sent over single TCP connection between client, server

$$(2RTT + t_0) \times 11 = \text{non-persistent no parallel}$$
$$(2RTT + t_0) \times 2 = \text{non-persistent parallel}$$

connections can be parallel

# Non-persistent HTTP

suppose user enters URL:

`www.someSchool.edu/someDepartment/home.index`

(contains text,  
references to 10  
jpeg images)

1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80

1b. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. "accepts" connection, notifying client

2. HTTP client sends HTTP request message (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`

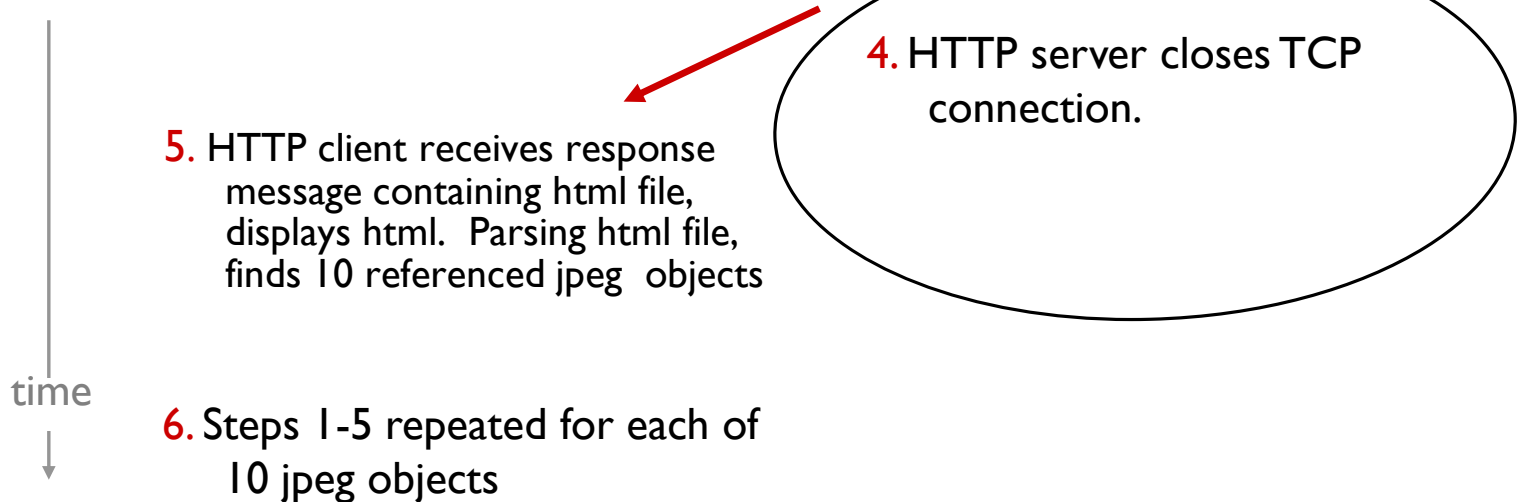
3. HTTP server receives request message, forms response message containing requested object, and sends message into its socket

time  
↓

RTT

to

# Non-persistent HTTP (cont.)






# Non-persistent HTTP: response time

**RTT (definition):** time for a small packet to travel from client to server and back

**HTTP response time:**

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time
- non-persistent HTTP response time =

$2\text{RTT} + \text{file transmission time}$



# Persistent HTTP

## *non-persistent HTTP issues:*

- requires 2 RTTs per object
- OS overhead for each TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

## *persistent HTTP:*

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects

$$RTT \times 12 + t_0 \times 11$$

# HTTP request message

- two types of HTTP messages: *request, response*
- **HTTP request message:**
  - ASCII (human-readable format)

request line  
(GET, POST,  
HEAD commands)

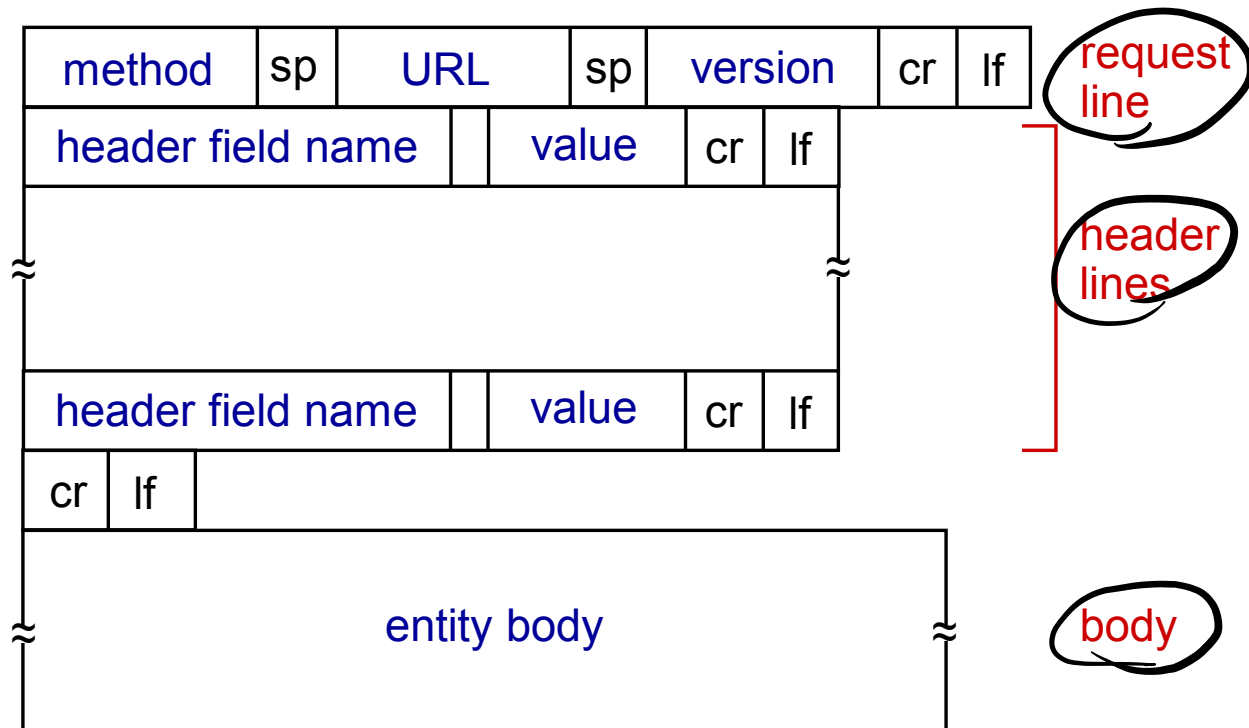
header  
lines

carriage return,  
line feed at start  
of line indicates  
end of header lines

```
GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Firefox/3.6.10\r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n
Keep-Alive: 115\r\n
Connection: keep-alive\r\n
\r\n
```

carriage return character  
line-feed character

# HTTP request message: general format



# Uploading form input

## POST method:

- web page often includes form input
- input is uploaded to server in entity body

## URL method:

- uses GET method
- input is uploaded in URL field of request line:

`www.somesite.com/animalsearch?monkeys&banana`

# Method types

## HTTP/1.0:

- GET
- POST
- HEAD
  - asks server to leave requested object out of response

## HTTP/1.1:

- GET, POST, HEAD
- PUT
  - uploads file in entity body to path specified in URL field
- DELETE
  - deletes file specified in the URL field

# HTTP response message

status line

(protocol  
status code  
status phrase)

header  
lines

data, e.g.,  
requested  
HTML file

```
HTTP/1.1 200 OK\r\n
Date: Sun, 26 Sep 2010 20:09:20 GMT\r\n
Server: Apache/2.0.52 (CentOS)\r\n
Last-Modified: Tue, 30 Oct 2007 17:00:02
GMT\r\n
ETag: "17dc6-a5c-bf716880"\r\n
Accept-Ranges: bytes\r\n
Content-Length: 2652\r\n
Keep-Alive: timeout=10, max=100\r\n
Connection: Keep-Alive\r\n
Content-Type: text/html; charset=ISO-8859-
1\r\n
\r\n
data data data data data ...
```

# HTTP response status codes

- status code appears in 1st line in server-to-client response message.

- some sample codes:

## **200 OK**

- request succeeded, requested object later in this msg

## **301 Moved Permanently**

- requested object moved, new location specified later in this msg (Location:)

## **400 Bad Request**

- request msg not understood by server

## **404 Not Found**

- requested document not found on this server

## **505 HTTP Version Not Supported**



# User-server state: cookies

many Web sites use cookies

*four components:*

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

**example:**

- Susan always access Internet from PC
- visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
  - unique ID
  - entry in backend database for ID

# Cookies: keeping “state” (cont.)

client



*maintain the cookie number  
for every server*



server



cookie file

usual http request msg

Amazon server  
creates ID  
1678 for user

usual http response  
set-cookie: 1678

create entry backend database




usual http request msg  
**cookie: 1678**

cookie-specific  
action

access

usual http response msg



access

cookie-specific  
action

one week later:



usual http request msg  
**cookie: 1678**

usual http response msg

# Cookies (continued)

*what cookies can be used for:*

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

*how to keep “state”:*

- protocol endpoints: maintain state at sender/receiver over multiple transactions
- cookies: http messages carry state

aside

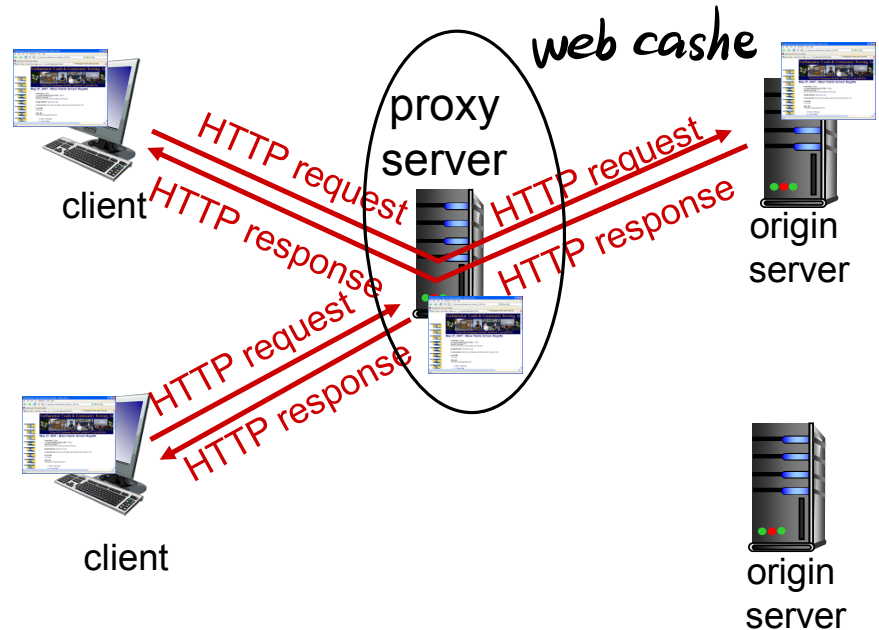
*cookies and privacy:*

- cookies permit sites to learn a lot about you
- you may supply name and e-mail to sites

# Web caches (proxy server)

*goal:* satisfy client request without involving origin server

- user sets browser: Web accesses via cache
- browser sends all HTTP requests to cache
  - object in cache: cache returns object
  - else cache requests object from origin server, then returns object to client



# More about Web caching

- cache acts as both client and server
  - server for original requesting client
  - client to origin server
- typically cache is installed by ISP (university, company, residential ISP)

## *why Web caching?*

- reduce response time for client request
- reduce traffic on an institution's access link
- Internet dense with caches: enables “poor” content providers to effectively deliver content (so too does P2P file sharing) ?

# Caching example:


## *assumptions:*

- avg object size: 1 M bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to all browsers: 15 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 15.4 Mbps

## *consequences:*

- LAN utilization:  $15\text{M}/100\text{M}=15\%$
- access link utilization =  $15/15.4=$  **99%**
- total delay = Internet delay + access delay + LAN delay  
= 2 sec + minutes + usecs

*problem!*




# Caching example: fatter access link

## *assumptions:*

- avg object size: 1 M bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 15 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: ~~15.4 Mbps~~ → 150 Mbps

## *consequences:*

- LAN utilization: 15%
- access link utilization = ~~99%~~ → 10%
- total delay = Internet delay + access delay + LAN delay  
= 2 sec + ~~minutes~~ → usecs



**Cost:** increased access link speed (not cheap!)

# Caching example: install local cache

## *assumptions:*


- avg object size: 1 M bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 15 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 15.4 Mbps

## *consequences:*

- LAN utilization: 15%
- access link utilization = ?
- total delay = ?

*How to compute link utilization, delay?*

*Cost:* web cache (cheap!)






# Caching example: install local cache

## *Calculating access link utilization, delay with cache:*

- suppose cache hit rate is 0.4
  - 40% requests satisfied at cache, 60% requests satisfied at origin
- access link utilization:
  - 60% of requests use access link
- data rate to browsers over access link  
 $= 0.6 * 15 \text{ Mbps} = 9 \text{ Mbps}$ 
  - utilization =  $9 / 15.4 = \underline{0.58}$
- total delay
  - ★  $= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
  - $= 0.6 (2.01) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$
  - less than with 150 Mbps link (and cheaper too!)



# Conditional GET

- **Goal:** don't send object if cache has up-to-date cached version

- no object transmission delay
- lower link utilization

- **cache:** specify date of cached copy in HTTP request

**If-modified-since:**  
**<date>**

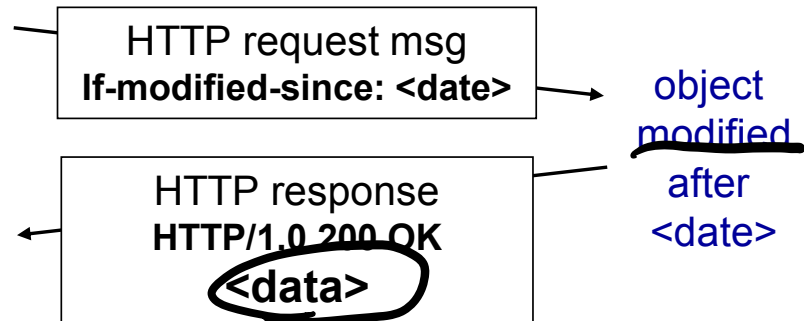
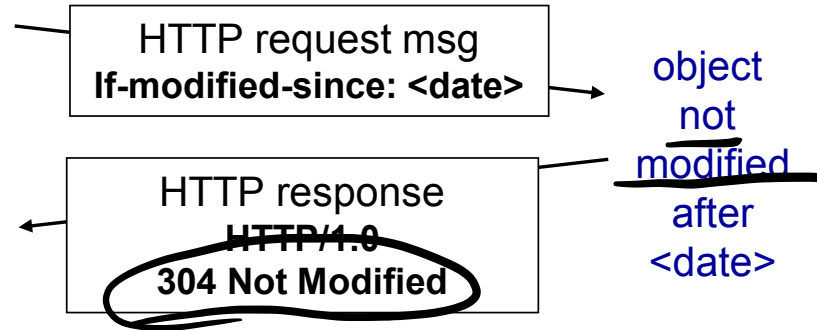
- **server:** response contains no object if cached copy is up-to-date:

**HTTP/1.0 304 Not Modified**

client



server



- ▲ `http`: Fetch html file first, parse it and post the request
- ▲ `request`: 一个 request 仅传回一个 object